

Rohan Siva

rohansiva@utexas.edu • (669) 294-7436 • [linkedin.com/in/rohan-siva](https://www.linkedin.com/in/rohan-siva) • github.com/Rohan-Siva • [Google Scholar](https://scholar.google.com/citations?user=rohansiva)

EDUCATION

University of Texas at Austin

Expected: May 2027

B.S. Electrical and Computer Engineering Honors, Minor in Stats & Data Science

GPA: 3.97/4.0

Relevant Coursework: Data Structures, Algorithms, Machine Learning, Data Science Principles, Data Science Lab, Probability and Statistics, Object-Oriented Design, Discrete Math, Multivariable Calculus, Linear Algebra, Differential Equations, Embedded Systems

EXPERIENCE

Decagon | Agent Software Engineer Intern | San Francisco, CA

June 2026 - Aug 2026

- Developing multimodal AI Voice Agents that manage complex customer interactions and drive customer retention

UT Austin Center for Autonomy | Undergraduate Research Intern | Austin, TX

Aug 2024 – Present

- Co-authored 6 research papers and presented work at top conferences on Autonomous Planning, Perception and Computer Vision
- Redesigned VLM training with distributed PyTorch and sharded dataloaders, achieving 3× faster fine-tuning and +15% accuracy
- Deployed VLM on Unitree Go2 via FastAPI microservice with load-balanced inference nodes, cutting inference latency by 38%

Mercor | Machine Learning Engineer Intern | Remote

Aug 2025 - Dec 2025

- Architected end-to-end LLM benchmarking pipeline for OpenAI MLE tasks using Retrieval-Augmented Generation (RAG) and PyTorch, benchmarking over 12,000+ OpenAI MLE tasks/month and reducing manual review time by 85%
- Scaled system to support 3× higher concurrent load (40 → 120 req/sec) with AWS Lambda concurrency, keeping <0.5% error rate
- Overhauled LLM evaluation pipeline, with sub-200ms latency and cut compute costs by 30% via optimized embedding caching

Cisco Systems | Software Engineer Intern | San Francisco, CA

May 2025 – Aug 2025

- Created an agent to convert natural language into ETL pipelines with direct enterprise PostgreSQL data loading, adopted by 9+ internal teams, decreasing pipeline-creation tickets by 80% per quarter and reducing setup time by 70%
- Built full-stack microservice with React, Elasticsearch, Bedrock and AWS Lambda for deployment, delivering 99.7% uptime
- Eliminated release failures with CI/CD for Docker and Kubernetes deployments and unit tests, achieving 90%+ coverage

Canyon Technologies | Software Engineer Intern | San Jose, CA

May 2024 – Aug 2024

- Used SpringBoot and Docker to build a cloud inventorying dashboard, serving 1,000+ users; Validated via Selenium for QA
- Integrated a secure Java and React.js mailing and password system with Firebase, integrating Google Single Sign-On (OAuth 2.0)

SELECTED PUBLICATIONS

Siva, R., Cheung, K., Li, L., & Sundaram, G. “kRAIG: A Natural Language-Driven Agent for Automated DataOps Pipeline Generation,” Under Review, International Conference on Machine Learning (ICML), 2026.

Bhatt, N. P., Yang, Y., Siva, R., Samineni, P., Milan, D., Wang, Z., & Topcu, U. “VLN-Zero: Rapid Exploration and Cache-Enabled Neurosymbolic Vision-Language Planning for Zero-Shot Transfer in Robot Navigation,” In Proceedings, NeuS, 2026. (Oral)

Bhatt, N. P., Yang, Y., Siva, R., Milan, D., Topcu, U., & Wang, Z. “Know Where You’re Uncertain When Planning with Multimodal Foundation Models: A Formal Framework,” In Proceedings, Conference on Machine Learning and Systems (MLSys), 2025. (Oral)

PROJECTS

MoodScribe: Emotion-based Chatbot | [Github](#)

Nov 2024

- Architected a full-stack web app with CyberPunk-themed UI (HTML/CSS) to generate sentiment-based Instagram reply suggestions, implementing async Meta API scraping, caching, and FastAPI endpoints to serve 2500+ concurrent users
- Fine-tuned Llama 3.1 using Unsloth and quantized LoRA adapters for sentiment analysis, cutting inference latency by 44%

NutriCheck: Image Classification for Nutrition | [Github](#)

Oct 2024

- Built an image classification app with custom Keras CNNs for real-time food analysis, retrieving nutritional data from Supabase
- Fine-tuned MobileNetV2 to classify food and display top 3 predictions with 92% accuracy, and output key nutritional content

Minerva: AI-Powered Study App | [Github](#)

Dec 2024

- Engineered a Flask and React AI study app for 300+ UT students with OpenAI API for tailored quizzes, flashcards and chatbot
- Developed a Retrieval-Augmented Generation pipeline using LlamaIndex to retrieve study material via Google Classroom API

SKILLS & INTERESTS

Languages: Python, Java, C/C++, JavaScript, SQL, TypeScript, HTML/CSS, Swift

Frameworks & Tools: React.js, Node.js, Flask, Docker, Kubernetes, PostgreSQL, AWS, Git, CI/CD, Azure, GCP, NumPy, Pandas

AI/ML: Pytorch, Keras, Sklearn, OpenCV, LangChain, Transformers, Apache Spark, Apache Kafka

Interests/Hobbies: Association for Computing Machinery (ACM) - Internal Officer, UT Convergent - AI Dev, VEX Robotics