

Rohan Siva

rohansiva@utexas.edu • (669) 294-7436 • [linkedin.com/in/rohan-siva](https://www.linkedin.com/in/rohan-siva) • github.com/Rohan-Siva • [Google Scholar](https://scholar.google.com/citations?user=rohansiva)

EDUCATION

University of Texas at Austin

Expected: May 2027

B.S. Electrical and Computer Engineering Honors, Minor in Stats & Data Science

GPA: 3.97/4.0

Relevant Coursework: Data Structures, Algorithms, Machine Learning, Data Science Principles, Data Science Lab, Probability and Statistics, Object-Oriented Design, Discrete Math, Multivariable Calculus, Linear Algebra, Differential Equations, Embedded Systems

SELECTED PUBLICATIONS

Siva, R., Bhatt, N. P., Yang, Y., Lee, S., Gadde, N., Ellis, C., Velasquez, A., Wang, Z., & Topcu, U. "What Objects Enable, Not What They Are: Functional Latent Spaces for Affordance Reasoning," Under Review, **Conference on Robot Learning (CoRL), 2026.**

Bhatt, N. P., Yang, Y., **Siva, R.,** Samineni, P., Milan, D., Wang, Z., & Topcu, U. "VLN-Zero: Rapid Exploration and Cache-Enabled Neurosymbolic Vision-Language Planning for Zero-Shot Transfer in Robot Navigation," In Proceedings, **NeuS, 2026. (Spotlight)**

Bhatt, N. P., Yang, Y., **Siva, R.,** Milan, D., Topcu, U., & Wang, Z. "Know Where You're Uncertain When Planning with Multimodal Foundation Models: A Formal Framework," In Proceedings, Conference on Machine Learning and Systems (**MLSys**), **2025. (Oral)**

EXPERIENCE

Decagon | *Research Intern* | San Francisco, CA

June 2026 - Present

- Post-trained open source voice models to add in-line pace, pause, and emotion control for a production agent with 99.5% stability
- Ran 90+ fine-tuning experiments, root-caused model architectural limitations, and designed reward functions for GRPO
- Shipped audio signal-to-noise ratio gate, filtering ~10% of noisy speech segments, and improving turn detection by 15%

UT Austin Center for Autonomy | *Researcher* | Austin, TX

Aug 2024 – Present

- Co-authored 6 research papers and presented work at top conferences on Autonomous Planning, Perception, and Computer Vision
- Redesigned VLM training with distributed PyTorch and sharded dataloaders, achieving 3× faster fine-tuning with +15% accuracy

Mercor | *Machine Learning Engineer Intern* | Remote

Aug 2025 - Dec 2025

- Architected end-to-end LLM benchmarking pipeline for OpenAI MLE tasks using Retrieval-Augmented Generation (RAG) and PyTorch, benchmarking over 12,000+ OpenAI MLE tasks/month and reducing manual review time by 85%
- Scaled system to support 3× higher concurrent load (40 → 120 req/sec) with AWS Lambda concurrency, keeping <0.5% error rate
- Overhauled LLM evaluation pipeline, with sub-200ms latency and cut compute costs by 30% via optimized embedding caching

Statistical Learning & AI Group Lab @ UT Austin | *Researcher*

May 2024 – Aug 2024

- Optimized speculative decoding for GRPO, increasing acceptance rate by 15% via training trajectory-conditioned EAGLE drafter
- Explored custom GRPO and PPO implementations for NanoChat and improved post-RL GSM8K performance by ~12%

Cisco Systems | *Machine Learning Engineer Intern* | San Francisco, CA

May 2025 – Aug 2025

- Created an agent to convert natural language into ETL pipelines with direct enterprise PostgreSQL data loading, adopted by 9+ internal teams, decreasing pipeline-creation tickets by 80% per quarter and reducing setup time by 70%. Published as a paper.
- Built full-stack microservice with React, Elasticsearch, Bedrock, and AWS Lambda for deployment, delivering 99.7% uptime
- Eliminated release failures with CI/CD for Docker and Kubernetes deployments and unit tests, achieving 90%+ coverage

PROJECTS

MoodScribe: Emotion-based Chatbot | [Github](#)

Nov 2024

- Architected a full-stack web app with CyberPunk-themed UI (HTML/CSS) to generate sentiment-based Instagram reply suggestions, implementing async Meta API scraping, caching, and FastAPI endpoints to serve 2500+ concurrent users
- Fine-tuned Llama 3.1 using Unsloth and quantized LoRA adapters for sentiment analysis, cutting inference latency by 44%

NutriCheck: Image Classification for Nutrition | [Github](#)

Oct 2024

- Built an image classification app with custom Keras CNNs for real-time food analysis, retrieving nutritional data from Supabase
- Fine-tuned MobileNetV2 to classify food and display top 3 predictions with 92% accuracy, and output key nutritional content

Minerva: AI-Powered Study App | [Github](#)

Dec 2024

- Engineered a Flask and React AI study app for 300+ UT students with OpenAI API for tailored quizzes, flashcards and chatbot
- Developed a Retrieval-Augmented Generation pipeline using LlamaIndex to retrieve study material via Google Classroom API

SKILLS & INTERESTS

Languages: Python, Java, C/C++, JavaScript, SQL, TypeScript, HTML/CSS, Swift

Frameworks & Tools: React.js, Node.js, Flask, Docker, Kubernetes, PostgreSQL, AWS, Git, CI/CD, Azure, GCP, NumPy, Pandas

AI/ML: Pytorch, Keras, Sklearn, OpenCV, LangChain, Transformers, Apache Spark, Apache Kafka